

Historical Training Data Manifest

This file describes the local training corpus without placing the corpus in the Git repository.

Dataset location

- Local path: `C:\Cprojectsgpu_app\MarketPinPredictor.worktrees\copilot\ensure-app-runs-correctly\historical_data`
- Repository policy: local-only; the `historical_data/` rule in `.gitignore` prevents accidental staging.
- Inventory source: recovery worktree on branch `friday-1/9`.
- Inventory date: 2026-09-18.

Warning: This location is temporary and data-bearing. The historical files must not be deleted, moved, or assumed to be disposable until a verified migration and training-readiness check is complete. Do not remove the old recovery worktree or its dataset until that migration has been independently verified.

Inventory

- Files: 10,227 CSV files and 2 pickle files.
- Size: approximately 114.64 GiB.
- Top-level groups: `aggregates/`, `combined_indices/`, `options/`, `s3_indices/`, and `training_datasets/`.
- File modification timestamps observed: 2026-09-16 00:52 through 2026-09-16 01:04 local time.
- Market-data coverage date range: unknown pending the full content inventory.

Use

The corpus is retained for historical feature generation, validation, and backtesting. The unclassified `historical_data/options/` subtree must not be used for training, validation, or backtesting until a full hash/content inventory confirms that it is not duplicated and is correctly symbol-filtered. Any

training job using a separately classified subset must document the dataset location and use chronological train, validation, and test splits to avoid future-data leakage.

Versioning rule

Git tracks the training code, schemas, configuration templates, and this manifest. It does not track the raw CSV corpus, generated features, model weights, caches, or runtime databases. A future dataset refresh should update this manifest with the inventory date, file count, size, date range, and a separate checksum manifest before older data is archived.